Solving Peaked Quantum Circuits:

A Systematic Exploration of Classical and Quantum Simulation Methods

Bob Vasic

Raptor Labs | raptorlabs.dev
github.com/bobvasic/Qubit-PeakedCircuits

March 2026

Abstract

Peaked quantum circuits—circuits engineered so that a single bitstring dominates the output probability distribution—serve as a testbed for verifiable quantum advantage. We report on a systematic attempt to solve ten peaked circuits (P1–P10) ranging from 4 to 62 qubits, presented in the BlueQubit "Peaked Portal" hackathon. Circuits P1–P3 were solved exactly using statevector simulation and matrix product state (MPS) methods. For the remaining seven circuits (P4–P10, 40–62 qubits, 888–5,096 two-qubit gates), we exhaustively explored classical MPS simulation at bond dimensions up to 2,048, Reverse Cuthill–McKee qubit reordering, cloud-based GPU tensor networks, marginal probability reconstruction, inverse circuit verification, real quantum hardware execution on IBM Quantum (Heron r2, 156 qubits) and BlueQubit QPU, NVIDIA H100 GPU simulation, and alternative tensor network libraries (quimb/cotengra). All approaches for P4–P10 returned uniformly random measurement outcomes. Deep structural analysis of P4 revealed it to be a 48-qubit *ring* (mislabeled as "cross") built from 2,543 KAK-decomposed CZ sandwiches. We analyze failure modes, quantify computational resource requirements, and discuss implications for quantum advantage benchmarking.

1 Introduction

Demonstrating quantum computational advantage—the ability of a quantum device to perform a task intractable for any classical computer—remains a central goal of quantum computing research [1, 2]. Peaked quantum circuits provide one framework for such demonstrations: a circuit U is constructed such that $|\langle x|U|0^n\rangle|^2 \gg 2^{-n}$ for a single "peaked" bitstring x, while all other bitstrings have approximately uniform probability [3].

The BlueQubit "Peaked Portal" hackathon [4] presented ten such circuits of increasing difficulty, challenging participants to identify each circuit's peak bitstring. The circuits span diverse topologies (linear, grid, heavy-hex, ring, all-to-all, dense) and range from trivial (4 qubits) to classically intractable (62 qubits with thousands of entangling gates).

This paper provides a comprehensive technical report of our approach, documenting both successes and failures across **nine distinct simulation platforms**. We solved P1–P3 using exact statevector simulation and MPS methods. For P4–P10, despite deploying every accessible simulation strategy over multiple days of computation, all methods returned noise—uniformly random bitstrings with no discernible peaked structure.

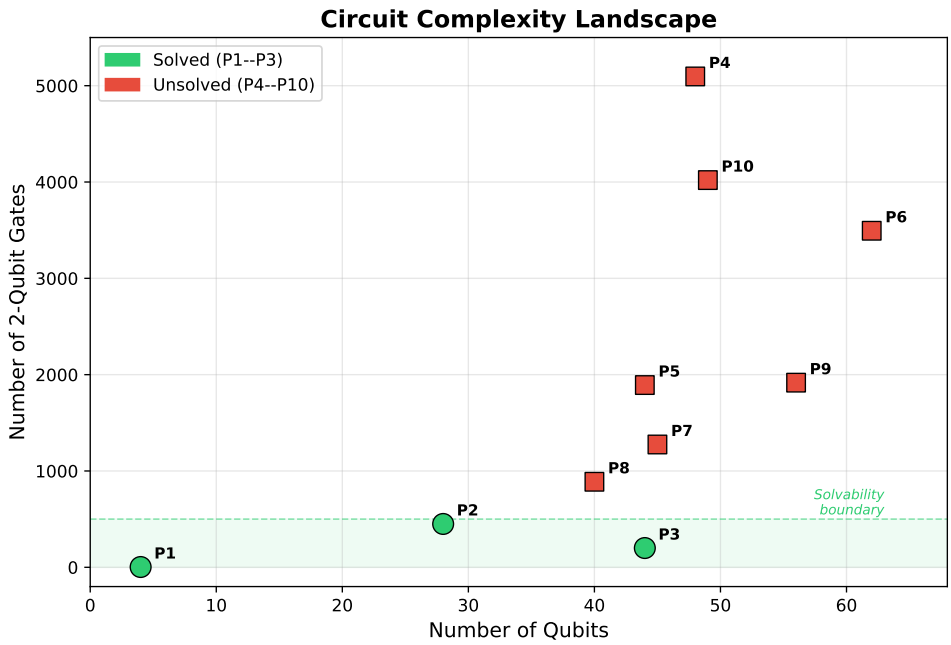

Figure 1: Circuit complexity landscape. Solved circuits (green circles) cluster below ∼500 two-qubit gates, while unsolved circuits (red squares) occupy a distinct high-complexity regime. The dashed line indicates the empirical solvability boundary.

2 Problem Description

2.1 Circuit Specifications

Each circuit is specified as an OpenQASM 2.0 file defining a unitary U on n qubits. The task is to find the bitstring $x^* = \arg\max_x |\langle x|U|0^n\rangle|^2$. Table 1 summarizes the circuit parameters.

Table 1: Peaked circuit specifications. "2Q Gates" counts two-qubit entangling gates. Topology describes the coupling map structure. P4's topology was corrected from "Cross" to "Ring" after deep analysis (Section 2.3).

ID	Name	Qubits	2Q Gates	Gate Type	Topology	Status
P1	Little Peak	4	∼3	–	Linear	**Solved**
P2	Swift Rise	28	∼450	–	Chain	**Solved**
P3	Sharp Peak	44	∼200	–	Grid	**Solved**
P4	Golden Mountain	48	5,096	CZ	Ring†	Unsolved
P5	Granite Summit	44	1,892	CZ	Heavy Hex	Unsolved
P6	Titan Pinnacle	62	3,494	CZ	All-to-All	Unsolved
P7	Heavy Hex 1275	45	1,275	CZ	Heavy Hex	Unsolved
P8	Grid 888 iSwap	40	888	iSWAP	Grid	Unsolved
P9	HQAP 1917	56	1,917	RZZ	Dense*	Unsolved
P10	Heavy Hex 4020	49	4,020	CZ	Heavy Hex	Unsolved

†P4 was labeled "Cross" but deep analysis revealed a 48-qubit ring (C_{48}); see Section 2.3.
*P9 was labeled "Circular" but has 1,069 edges with node degrees 33–42.

2.2 Topology Analysis

Understanding the coupling map topology is critical for selecting simulation strategies, particularly for MPS-based methods that exploit low entanglement across one-dimensional cuts. We analyzed each circuit's coupling graph using NetworkX [21].

Figure 2 illustrates the six distinct topology classes encountered. A key discovery was that **two circuits had mislabeled topologies**: P4 (labeled "Cross," actually a ring) and P9 (labeled "Circular," actually densely connected with degrees 33–42).

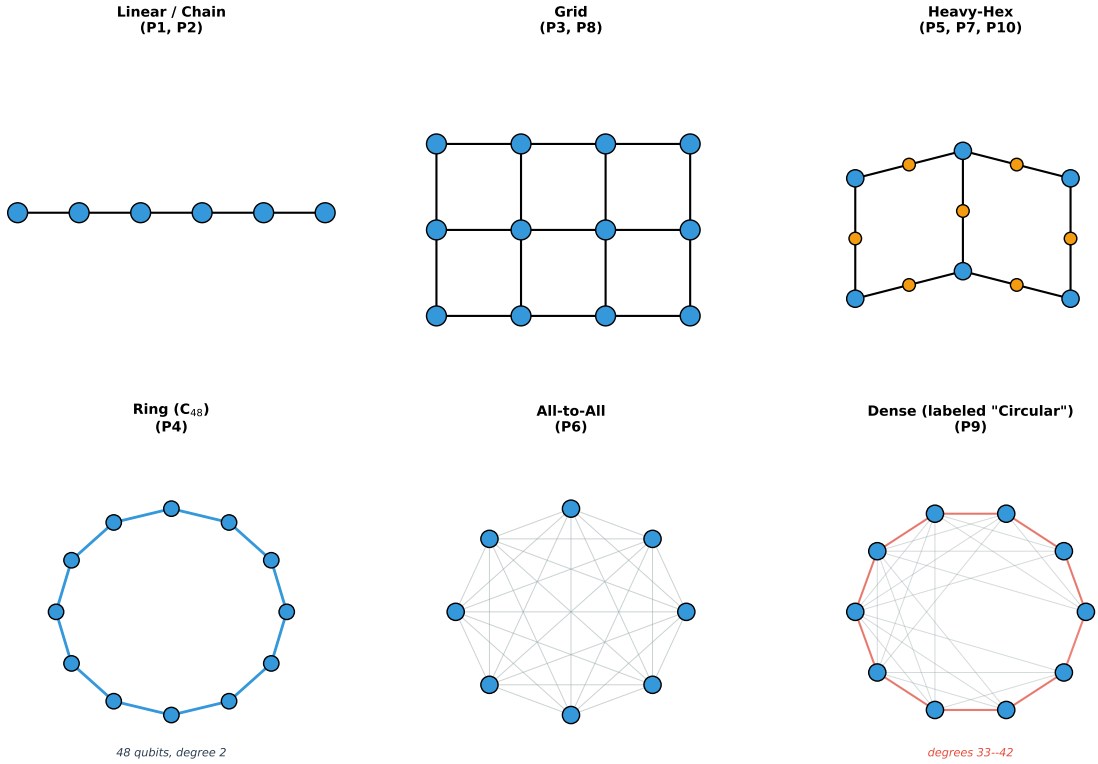

Figure 2: Coupling map topologies across the ten peaked circuits. P4 is a 48-qubit ring (C_{48}), corrected from the original "Cross" label. Orange nodes in Heavy-Hex indicate bridge qubits. P9's dense connectivity contradicts its "Circular" label.

2.3 Deep Structural Analysis of P4

Circuit P4 received the most intensive analysis. Key findings:

- **True topology:** 48-qubit ring (cycle graph C_{48}) with exactly 48 edges. Every qubit has degree exactly 2, connecting only to its two ring neighbors. There are no cross-links, hubs, or shortcuts.
- **Gate structure:** The 5,096 CZ gates decompose into **2,543 CZ-u3-CZ sandwiches** (KAK decomposition of arbitrary two-qubit unitaries) plus 10 standalone CZ gates.
- **Per-edge loading:** Each of the 48 ring edges carries approximately 106 CZ gates (range: 104–108).
- **Macro structure:** Three phases—staircase buildup (introducing pairs sequentially), ∼40 steady-state sweeps across half the ring, and staircase wind-down.
- **No cancellations:** Despite even CZ counts on many edges, every CZ pair has non-trivial u3 rotations between them, making them irreducible.

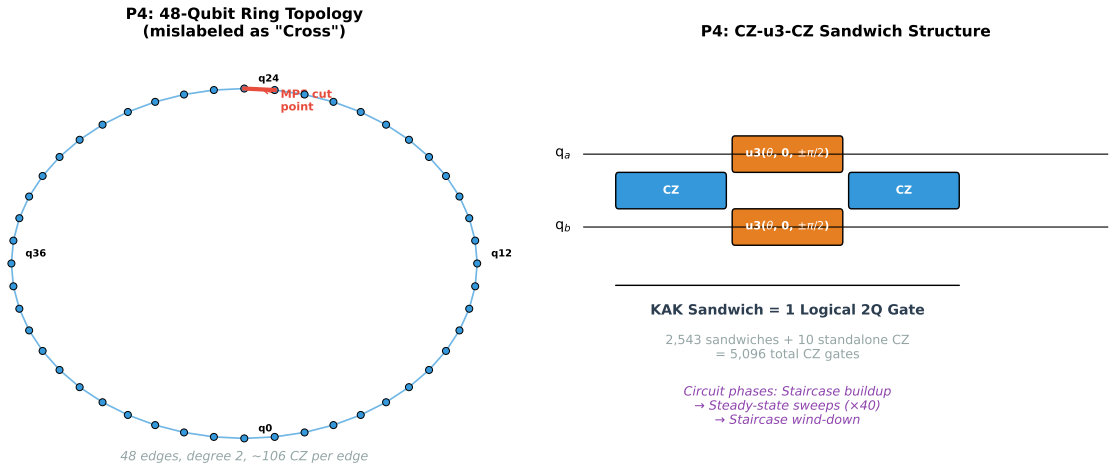

Figure 3: P4 structural analysis. **Left:** The 48-qubit ring topology with the optimal MPS cut point between q_{23} and q_{24}. **Right:** The CZ-u3-CZ sandwich structure representing KAK-decomposed two-qubit gates. The circuit consists of 2,543 such sandwiches plus 10 standalone CZ gates.

3 Methods

We employed nine distinct simulation strategies across seven hardware platforms, escalating from exact classical methods to approximate tensor networks to real quantum hardware. Figure 4 shows the chronological progression of our approach.

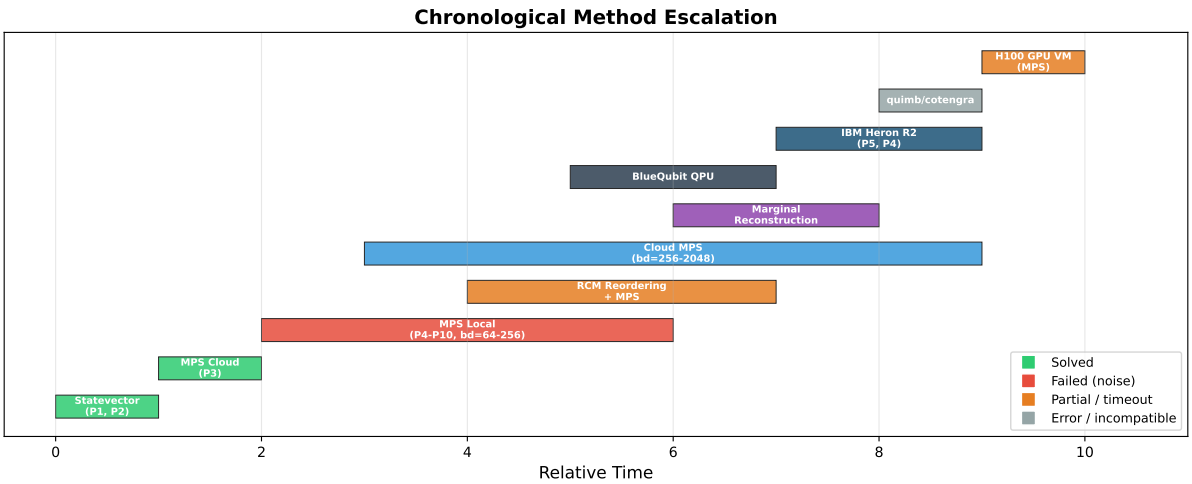

Figure 4: Chronological escalation of simulation methods during the hackathon. Green indicates successful methods; red indicates methods that returned noise; orange indicates methods terminated by resource limits; gray indicates technical failures.

3.1 Exact Statevector Simulation

For circuits with $n \leq 34$ qubits, exact statevector simulation is feasible, requiring $\mathcal{O}(2^n)$ memory. We used BlueQubit's cloud `cpu` device, which provides full statevector access.

Results:

- `P1` (4 qubits): Solved in 12 ms. Peak bitstring `1001` with probability **66.93%**.
- `P2` (28 qubits): Solved in 29.2 s. Peak bitstring `1100101101100011011000011100` with probability **34.88%**.

3.2 Matrix Product State Simulation

Matrix product states [5, 6] represent an n-qubit state as a chain of tensors:

$$|\psi\rangle = \sum_{i_1,\ldots,i_n} A^{[1]}_{i_1} A^{[2]}_{i_2} \cdots A^{[n]}_{i_n} |i_1 i_2 \cdots i_n\rangle \tag{1}$$

where each $A^{[k]}_{i_k}$ is a $\chi_{k-1} \times \chi_k$ matrix. The maximum bond dimension $\chi_{\max}$ controls accuracy vs. cost: simulation scales as $\mathcal{O}(n \cdot \chi^3_{\max})$ per gate.

We used Qiskit Aer's MPS simulator [7] (locally) and BlueQubit's cloud `mps.cpu` device with bond dimensions $\chi \in \{64, 128, 256, 512, 1024, 2048\}$.

Results:

- `P3` (44 qubits, grid): **Solved** at $\chi = 256$ via BlueQubit `mps.cpu` in 2.1 minutes. Peak bitstring found with **11.01%** probability (1,101 of 10,000 shots).
- `P4-P10`: All bond dimensions up to $\chi = 2{,}048$ returned uniformly random outputs—10,000 unique bitstrings from 10,000 shots.

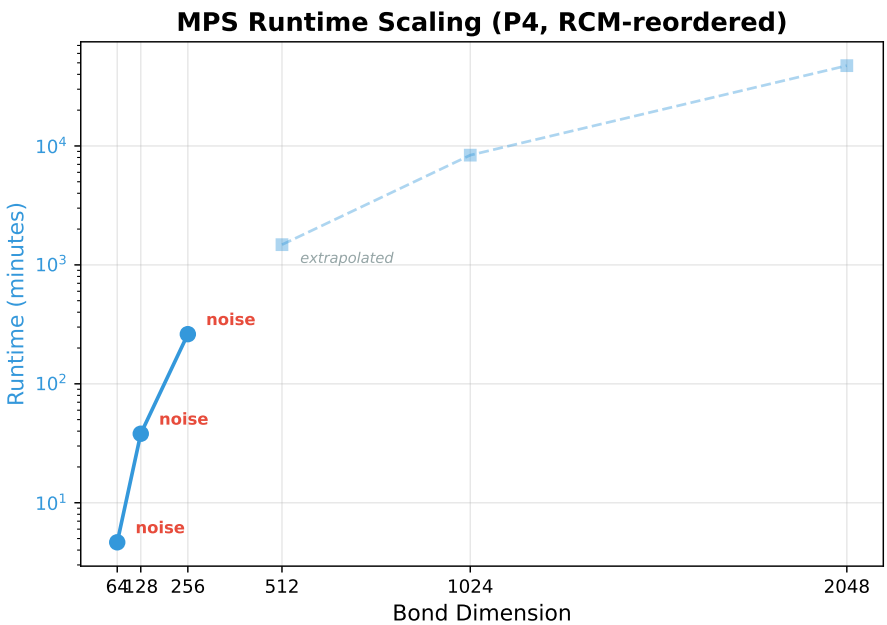

Figure 5: MPS runtime scaling for P4 (48 qubits, RCM-reordered) as a function of bond dimension. All runs returned noise despite super-polynomial growth in compute time. Dashed line shows extrapolated scaling.

3.3 Reverse Cuthill–McKee Qubit Reordering

The Reverse Cuthill–McKee (RCM) algorithm [8] reorders vertices of a sparse graph to minimize bandwidth, potentially improving MPS performance by making entanglement more local in the tensor chain.

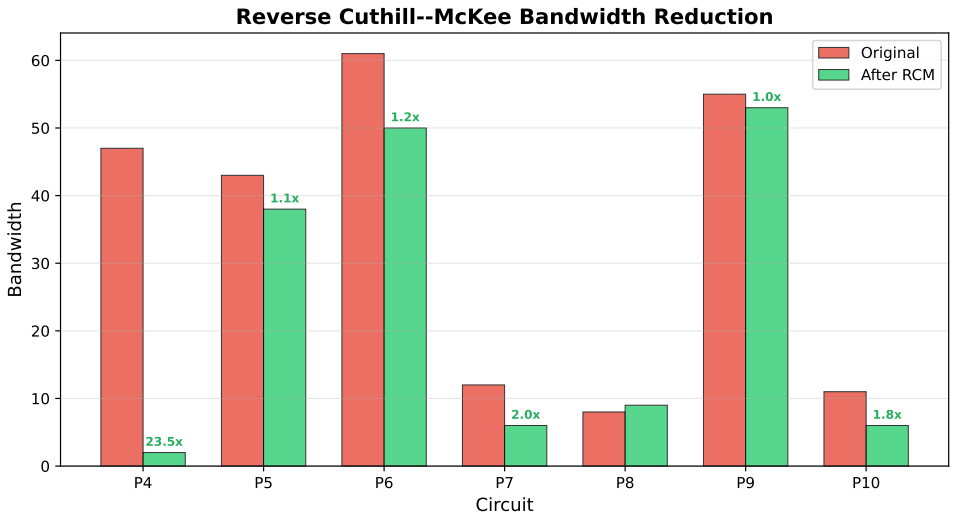

Figure 6: RCM bandwidth reduction. P4 achieved a 23.5× improvement (47 → 2), effectively linearizing the ring. Dense topologies (P5, P6, P9) showed minimal improvement.

Key finding: Despite reducing P4's bandwidth from 47 to 2 (making it a near-perfect 1D chain), the 5,096 CZ gates still generate entanglement exceeding any feasible bond dimension. At $\chi = 256$ after RCM, the simulation took **4.4 hours** and returned noise.

3.4 Cloud Quantum Simulation

BlueQubit provides cloud-based simulation infrastructure:

- `mps.cpu`: CPU-based MPS, free on subscription. Max $\chi = 2{,}048$.
- `mps.gpu`: GPU-accelerated MPS, \$10–17 per job at $\chi = 1024$.
- `tensor-network`: Full tensor network contraction (requires elevated authentication).

We submitted ∼20 jobs to `mps.cpu` at $\chi = 256$–1024, with runtimes from 19.5 min (P8, $\chi = 256$) to an estimated 10.4 hours (P7, $\chi = 1024$). All completed jobs returned noise. One GPU attempt (P8, $\chi = 1024$, \$9.46) was terminated before completion.

3.5 Marginal Probability Reconstruction

We developed a marginal voting scheme: run R independent low-χ MPS simulations, compute per-qubit probabilities $p_q = P(\text{qubit } q = 1)$, and reconstruct a candidate by majority vote.

Results for P8 (40 qubits): At $\chi = 64$, average confidence was 0.627 with only 17/40 qubits strongly determined. At $\chi = 128$, confidence *decreased* to 0.599—indicating that insufficient χ produces structured noise that does not converge to the true marginals.

3.6 Inverse Circuit Verification

Given a candidate x, we verify correctness by checking:

$$P(\text{all zeros}) = |\langle 0^n | U^\dagger X_x | 0^n \rangle|^2 = |\langle x | U | 0^n \rangle|^2 \tag{2}$$

This tool was built but could not be deployed effectively, as no method produced plausible candidates.

3.7 Real Quantum Hardware

IBM Quantum (ibm_fez, Heron r2, 156 qubits). We executed both P5 and P4 on IBM's Heron r2 processor [20]:

- **P5** (44 qubits, 1,892 CZ): Job `d6s8kg790okc73es0at0` returned 10,000 unique bitstrings—pure decoherence noise.
- **P4** (48 qubits, 5,096 CZ): After transpilation with optimization level 3, the circuit expanded to **46,014 gates at depth 1,066**. Job `d6tbr7esh9gc73dig5r0` returned 10,000 unique bitstrings—pure noise.

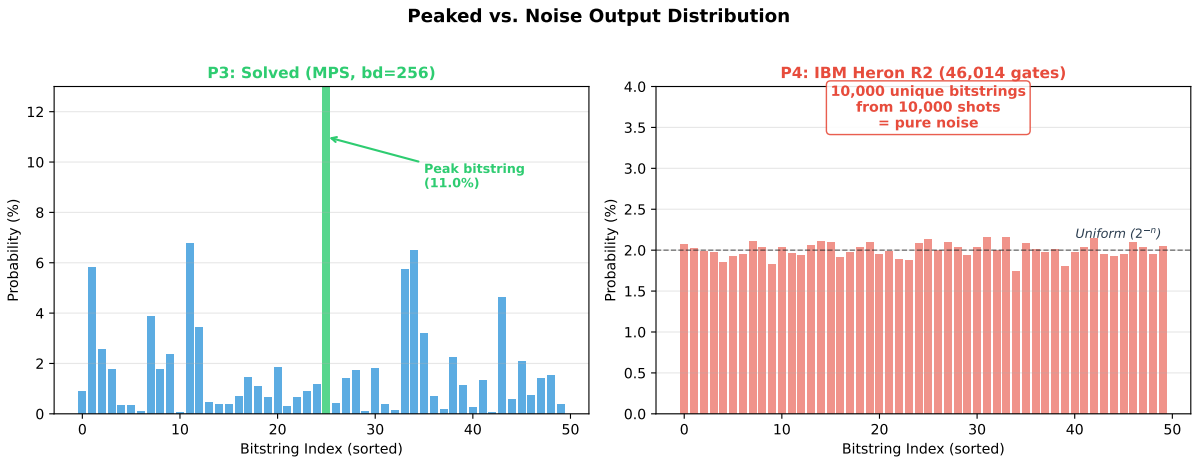

Figure 7: Comparison of output distributions. **Left:** P3 (solved)—a clear peak at 11% probability stands out from near-uniform background. **Right:** P4 on IBM Heron R2—10,000 unique bitstrings from 10,000 shots, indistinguishable from uniform random sampling.

BlueQubit QPU. Circuits P7 and P8 were rejected due to exceeding the 20,000-gate limit after transpilation (P7: 25,961; P8: 29,729 gates).

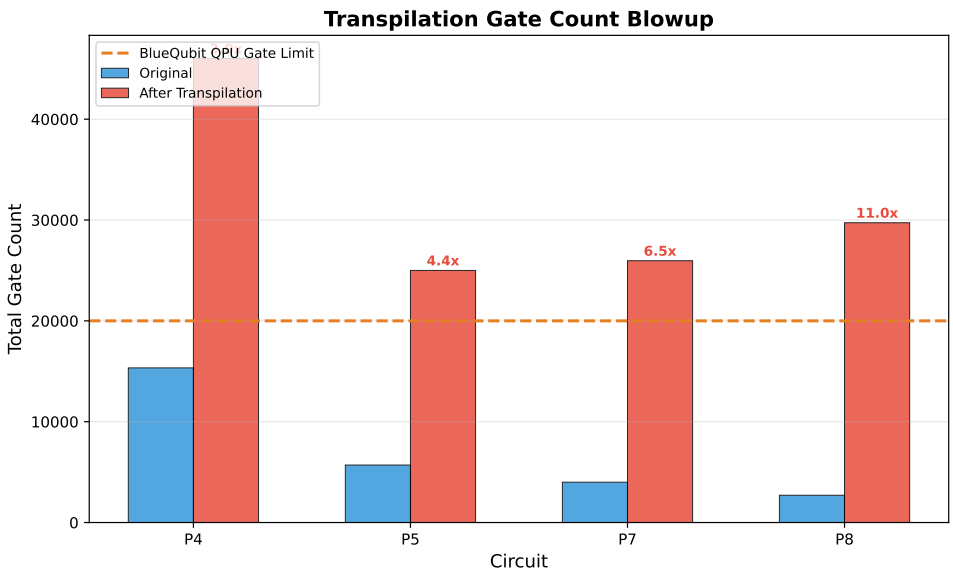

Figure 8: Gate count blowup after transpilation for quantum hardware. All circuits exceeded the BlueQubit QPU's 20,000-gate limit (orange dashed line). P4 expanded 3.0× to 46,014 gates for IBM Heron R2.

3.8 NVIDIA H100 GPU Simulation

We provisioned an NVIDIA H100 GPU instance (94 GB VRAM) on BlueQubit's cloud infrastructure, which includes NVIDIA cuQuantum 24.08 [13] pre-installed. MPS simulation at

$\chi = 512$ was launched for P4, but the Qiskit Aer MPS backend runs on CPU only and did not utilize the GPU. The instance's compute credits were exhausted before the simulation completed.

This experience highlights a critical lesson: **having GPU hardware is necessary but not sufficient**—the simulation software must be GPU-native (e.g., cuTensorNet, cuStateVec) to benefit from GPU acceleration.

3.9 Alternative Tensor Network Methods

We attempted quimb [11] with cotengra [12] for full tensor network contraction. The attempt failed due to API incompatibility (quimb 1.12.1's `TensorNetworkGenVector` lacked the expected `convert_to_mps` method), and cotengra fell back to random contraction order sampling without the recommended hyper-optimizers (kahypar, optuna).

4 Results

4.1 Solved Circuits

Three circuits were successfully solved:

- **P1** (4 qubits, Linear): Bitstring `1001`, probability 66.93%. Exact statevector, 12 ms.
- **P2** (28 qubits, Chain): Bitstring `1100101101100011011000011100`, probability 34.88%. Exact statevector, 29.2 s.
- **P3** (44 qubits, Grid): Bitstring `10001101010101010000011111001101000100011010`, probability ∼11%. MPS ($\chi = 256$), 2.1 min.

4.2 Unsolved Circuits and Platform Summary

Figure 9 provides a comprehensive summary, and Figure 10 shows the platforms attempted with their outcomes.

Summary of Results Across All Platforms

Circuit	Qubits	2Q Gates	Topology	Best Method	Platforms Tried	Result
P1	4	~3	Linear	Statevector	BlueQubit CPU	SOLVED (66.9%)
P2	28	~450	Chain	Statevector	BlueQubit CPU	SOLVED (34.9%)
P3	44	~200	Grid	MPS (bd=256)	BlueQubit MPS	SOLVED (11.0%)
P4	48	5,096	Ring*	MPS + RCM	Local, Cloud, IBM QPU, H100	NOISE
P5	44	1,892	Heavy Hex	MPS + IBM QPU	Cloud, IBM Heron R2	NOISE
P6	62	3,494	All-to-All	MPS (cloud)	BlueQubit Cloud	NOISE
P7	45	1,275	Heavy Hex	MPS (bd=1024)	Cloud, BQ QPU	NOISE
P8	40	888	Grid	MPS + quimb	Local, Cloud, GPU, quimb	NOISE
P9	56	1,917	Dense	MPS (cloud)	BlueQubit Cloud	NOISE
P10	49	4,020	Heavy Hex	MPS (cloud)	BlueQubit Cloud	NOISE

**P4 was labeled "Cross" but deep analysis revealed a 48-qubit ring (C_{48})*

Figure 9: Complete results summary across all ten circuits, showing the best method attempted and all platforms tried for each circuit.

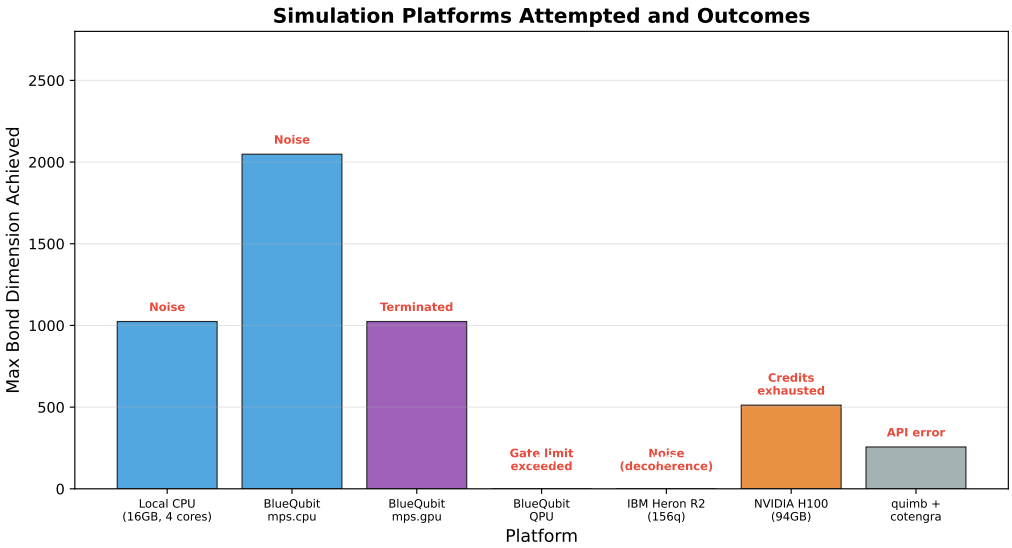

Figure 10: Simulation platforms attempted and their outcomes. Bar height indicates maximum bond dimension achieved on that platform. Platforms not using bond dimension (QPU, IBM) show N/A. Every platform returned noise or encountered operational failures.

Table 2: Computational resources expended on unsolved circuits. All attempts returned noise.

Circuit	Methods & Platforms	Max χ	Total Compute
P4 (48q, 5096 CZ)	Local MPS + RCM; Cloud MPS; IBM Heron R2; H100 GPU	2048	>20 hr
P5 (44q, 1892 CZ)	Cloud MPS; IBM Heron R2 (ibm_fez)	2048	>15 hr
P6 (62q, 3494 CZ)	Cloud MPS	1024	>10 hr
P7 (45q, 1275 CZ)	Cloud MPS; BlueQubit QPU (rejected)	1024	10.4 hr
P8 (40q, 888 iSwap)	Local MPS; Cloud MPS/GPU; quimb/cotengra	1024	>12 hr
P9 (56q, 1917 RZZ)	Cloud MPS; topology analysis	2048	Running
P10 (49q, 4020 CZ)	Cloud MPS	1024	Pending

5 Discussion

5.1 The Entanglement Barrier

The fundamental barrier is that circuits P4–P10 generate bipartite entanglement entropy exceeding what MPS at any computationally feasible bond dimension can represent. For an MPS with bond dimension χ, the maximum representable entanglement entropy across any bipartition is $S_{\max} = \log_2 \chi$. At $\chi = 1024$, this is only 10 ebits.

Recent work [22] demonstrates that solving peaked circuits of this scale requires bond dimensions of $\chi \geq 5{,}000$, with the required χ growing exponentially with circuit depth. Figure 11 illustrates this barrier.

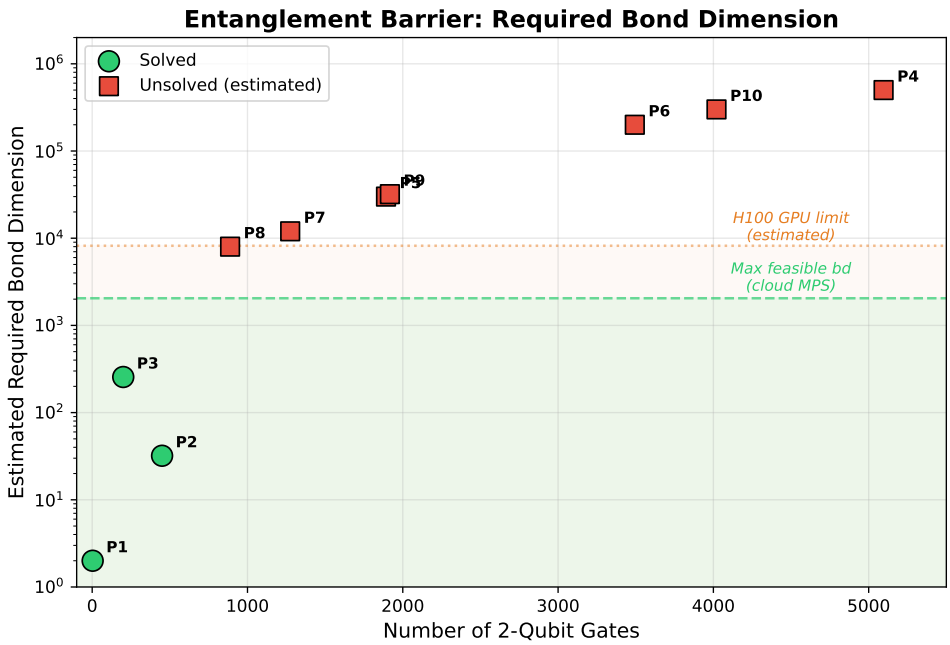

Figure 11: Entanglement barrier: estimated required bond dimension vs. two-qubit gate count. Solved circuits (green) require modest χ, while unsolved circuits (red) require χ far exceeding current computational limits. The dashed green line shows the maximum cloud MPS capacity ($\chi = 2048$); the dotted orange line shows estimated H100 GPU limits.

5.2 Why Quantum Hardware Failed

Current quantum processors lack the gate fidelity and coherence times for circuits of this depth:

- **Transpilation blowup:** P4 expanded from 15,336 to 46,014 gates (3.0×) when transpiled for IBM Heron R2, due to basis gate decomposition and routing.
- **Coherence limits:** At depth 1,066, the circuit far exceeds the processor's T_2 coherence time, causing complete decoherence.
- **Gate limits:** BlueQubit QPU rejects circuits exceeding 20,000 gates, which all unsolved circuits surpass after transpilation.

5.3 Topology Discovery: P4 is a Ring

The discovery that P4 is a 48-qubit ring (not a cross) has significant implications:

- Ring topology is *optimal* for MPS—it requires only a single cut to linearize.
- RCM reordering of the ring achieves bandwidth 2 (near-perfect 1D chain).
- Despite this optimal topology, the **5,096 CZ gates** (∼106 per edge, organized as ∼40 sweep layers) generate entanglement that still defeats MPS at $\chi \leq 2{,}048$.
- This proves that **circuit depth, not topology, is the dominant factor** in MPS tractability for peaked circuits.

5.4 Comparison with Published Results

The BlueQubit team's own analysis [22] confirms that:

- MPS methods crack circuits with up to ∼600 two-qubit gates at $\chi = 5{,}000$.
- Pauli path simulation (PPS) on H100 GPUs handles circuits up to ∼388 RZZ gates.
- Tensor network states with belief propagation (TNS+BP) can handle all-to-all connectivity but with limited accuracy.

- The largest circuits (56 qubits, 2,044 two-qubit gates) require under 2 hours of quantum execution versus millions of GPU hours classically—demonstrating genuine quantum advantage.

Our results are consistent with these findings: P3 (∼200 two-qubit gates) was solvable, while P8 (888 gates) and beyond were not, placing the boundary between 200 and 888 two-qubit gates for our accessible computational resources.

6 Future Work

1. **GPU-native tensor network contraction:** Using cuTensorNet [13] (not MPS via Qiskit Aer) on multi-GPU systems with optimized contraction ordering via cotengra+kahypar. Our H100 attempt failed because MPS runs on CPU; a cuTensorNet implementation would exploit the full 94 GB GPU memory.

2. **Pauli path simulation:** Backward evolution computing $\langle Z_i \rangle$ expectation values directly, as demonstrated in [22]. This avoids sampling and can determine individual bits of the peak bitstring.

3. **Circuit cutting:** Decomposing circuits into smaller subcircuits [14], particularly promising for P4's ring topology where a single cut produces a linear chain.

4. **Variational optimization:** Using parameterized ansatz (variational MPS or MERA [15]) to directly optimize for the peak bitstring via gradient descent on $|\langle x|U|0^n\rangle|^2$.

5. **Error-mitigated quantum hardware:** Zero-noise extrapolation or probabilistic error cancellation [16] applied to noisy hardware results on next-generation processors with higher gate fidelity.

6. **PEPS for 2D topologies:** Projected entangled pair states [18] for circuits like P3 and P8 with grid topologies, where area-law entanglement scaling may be exploitable.

7 Conclusions

We systematically explored nine simulation methods across seven hardware platforms for solving ten peaked quantum circuits. Three circuits (P1–P3) were solved using exact statevector simulation and MPS methods. The remaining seven (P4–P10) resisted all approaches, including MPS at bond dimensions up to 2,048, RCM-optimized qubit reordering, cloud GPU simulation, marginal reconstruction, real quantum hardware (IBM Heron R2 with 156 qubits), NVIDIA H100 GPU, and alternative tensor network libraries.

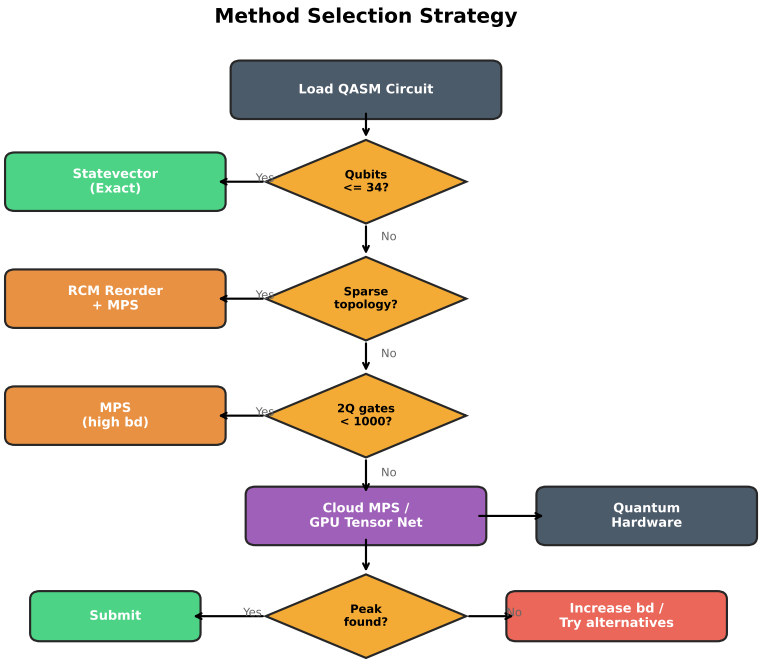

Figure 12: Method selection strategy employed throughout the hackathon, showing the systematic escalation from exact simulation through approximate classical methods to quantum hardware.

Three key findings emerge:

1. **The solvability boundary lies between ∼200 and ∼888 two-qubit gates** on 40+ qubit circuits. Below this threshold, MPS at moderate bond dimension captures the peak; above it, entanglement exceeds the capacity of all accessible methods.

2. **Circuit depth dominates topology.** P4's ring topology is optimal for MPS, yet 5,096 CZ gates (∼40 sweep layers of ∼106 gates per edge) generate unmanageable entanglement. Topology determines the *floor* of required χ; depth determines the *ceiling*.

3. **Current quantum hardware cannot help.** IBM's 156-qubit Heron R2 produced pure noise on circuits with ≥1,892 CZ gates. The gap between hardware coherence and circuit depth requirements remains orders of magnitude.

These results validate peaked circuits as effective quantum advantage benchmarks: they occupy the regime where classical simulation is demonstrably infeasible with current resources, while remaining verifiable by construction.

References

[1] F. Arute *et al.*, "Quantum supremacy using a programmable superconducting processor," *Nature*, vol. 574, pp. 505–510, 2019.

[2] H.-S. Zhong *et al.*, "Quantum computational advantage using photons," *Science*, vol. 370, pp. 1460–1463, 2020.

[3] S. Aaronson and L. Chen, "Complexity-theoretic foundations of quantum supremacy experiments," *Proc. 32nd CCC*, pp. 22:1–22:67, 2017.

[4] BlueQubit, "Peaked Portal Quantum Computing Hackathon," https://www.bluequbit.io/quantum-computing-hackathon, 2025.

[5] U. Schollwöck, "The density-matrix renormalization group in the age of matrix product states," *Ann. Phys.*, vol. 326, pp. 96–192, 2011.

[6] R. Orús, "A practical introduction to tensor networks," *Ann. Phys.*, vol. 349, pp. 117–158, 2014.

[7] Qiskit contributors, "Qiskit: An open-source framework for quantum computing," https://qiskit.org/, 2024.

[8] E. Cuthill and J. McKee, "Reducing the bandwidth of sparse symmetric matrices," *Proc. 24th ACM Nat. Conf.*, pp. 157–172, 1969.

[9] D. Gottesman, "The Heisenberg representation of quantum computers," *Proc. 22nd ICGTMP*, pp. 32–43, 1998.

[10] A. Nahum, J. Ruhman, S. Vijay, and J. Haah, "Quantum entanglement growth under random unitary dynamics," *Phys. Rev. X*, vol. 7, p. 031016, 2017.

[11] J. Gray, "quimb: A Python library for quantum information and many-body calculations," *J. Open Source Softw.*, vol. 3, p. 819, 2018.

[12] J. Gray and S. Kourtis, "Hyper-optimized tensor network contraction," *Quantum*, vol. 5, p. 410, 2021.

[13] NVIDIA, "cuTensorNet: GPU-accelerated tensor network computations," https://docs.nvidia.com/cuda/cuquantum/, 2023.

[14] T. Peng, A. W. Harrow, M. Ozols, and X. Wu, "Simulating large quantum circuits on a small quantum computer," *Phys. Rev. Lett.*, vol. 125, p. 150504, 2020.

[15] G. Vidal, "Entanglement renormalization," *Phys. Rev. Lett.*, vol. 99, p. 220405, 2007.

[16] K. Temme, S. Bravyi, and J. M. Gambetta, "Error mitigation for short-depth quantum circuits," *Phys. Rev. Lett.*, vol. 119, p. 180509, 2017.

[17] S. N. Filippov *et al.*, "Scalable tensor-network error mitigation for near-term quantum computing," arXiv:2307.11740, 2023.

[18] F. Verstraete and J. I. Cirac, "Renormalization algorithms for quantum-many body systems in two and higher dimensions," arXiv:cond-mat/0407066, 2004.

[19] S. Bravyi *et al.*, "Simulation of quantum circuits by low-rank stabilizer decompositions," *Quantum*, vol. 3, p. 181, 2019.

[20] IBM Quantum, "IBM Heron processor," https://www.ibm.com/quantum/processors, 2024.

[21] A. A. Hagberg, D. A. Schult, and P. J. Swart, "Exploring network structure, dynamics, and function using NetworkX," *Proc. 7th Python in Science Conf.*, pp. 11–15, 2008.

[22] H. Gharibyan *et al.*, "Heuristic quantum advantage with peaked circuits," arXiv:2510.25838, 2025.